
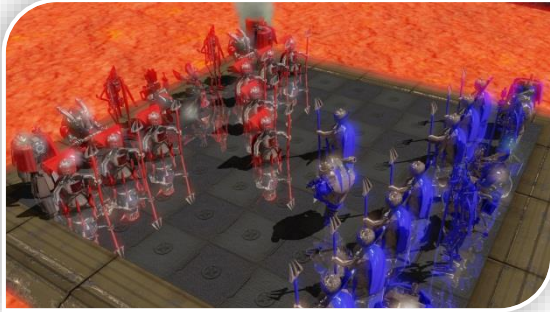


Game Engine Design



tum.3D
computer graphics & visualization

- This week: Explosions!



Task overview




1. Extend your sprite renderer to support animated sprites
2. Spawn and update explosions
3. Create an additional particle system

- Sequence of keyframes instead of a single, static image



- Internally: 2D texture array instead of 2D texture
 - Texture2DArray != Array of Texture2D
 - Extend sprite vertices by parameters $t, \alpha \in [0;1]$
 - t indicates animation progress of the sprite
 - α is a scalar opacity factor for fading sprites in and out

- Each texture resource in D3D11 can consist of multiple subresources



- The **subresource index** can be calculated with `D3D11CalcSubresource()`
- Subresources can be initialized individually directly in `ID3D11Device::CreateTexture2D()`

```
HRESULT ID3D11Device::CreateTexture2D(  
    const D3D11_TEXTURE2D_DESC *pDesc,           //description of texture  
    const D3D11_SUBRESOURCE_DATA *pInitialData, //pointer to an array of subresources  
                                                    //with pDesc->MipLevels*pDesc->ArraySize elements  
    ID3D11Texture2D **ppTexture2D               //output  
);
```

- pDesc->ArraySize: Size of array
- D3D11_SUBRESOURCE_DATA
 - Basically stores a pointer to subresource bytedata for initialization

```
typedef struct D3D11_SUBRESOURCE_DATA {  
    const void *pSysMem; //pointer to initialization data  
    UINT SysMemPitch;     //bytesize of a line in a 2D/3D texture  
    UINT SysMemSlicePitch; //bytesize of a xy-slice in a 3D texture  
} D3D11_SUBRESOURCE_DATA;
```

- Description for shader resource view needs to be set manually
 - ID3D11Device::CreateShaderResourceView(..., const D3D11_SHADER_RESOURCE_VIEW_DESC *pDesc, ...)

C++

```
typedef struct D3D11_SHADER_RESOURCE_VIEW_DESC {  
    DXGI_FORMAT          Format;  
    D3D11_SRV_DIMENSION ViewDimension;  
    union {  
        D3D11_BUFFER_SRV      Buffer;  
        D3D11_TEX1D_SRV       Texture1D;  
        D3D11_TEX1D_ARRAY_SRV Texture1DArray;  
        D3D11_TEX2D_SRV       Texture2D;  
        D3D11_TEX2D_ARRAY_SRV Texture2DArray;  
        D3D11_TEX2DMS_SRV     Texture2DMS;  
        D3D11_TEX2DMS_ARRAY_SRV Texture2DMSArray;  
        D3D11_TEX3D_SRV       Texture3D;  
        D3D11_TEXCUBE_SRV     TextureCube;  
        D3D11_TEXCUBE_ARRAY_SRV TextureCubeArray;  
        D3D11_BUFFEREX_SRV    BufferEx;  
    };  
} D3D11_SHADER_RESOURCE_VIEW_DESC;
```


- Format = „texture format“ (DXGI_FORMAT_*, from texture description)
- ViewDimension = D3D11_SRV_DIMENSION_TEXTURE2DARRAY
- Texture2DArray:
 - ArraySize = „size of array“ (from texture description)
 - MipLevels = „number of miplevels“ (from texture description)
 - FirstArraySlice = MostDetailedMip = 0

C++

```
typedef struct D3D11_SHADER_RESOURCE_VIEW_DESC {  
    DXGI_FORMAT      Format;  
    D3D11_SRV_DIMENSION ViewDimension;  
    union {  
        D3D11_BUFFER_SRV      Buffer;  
        D3D11_TEX1D_SRV      Texture1D;  
        D3D11_TEX1D_ARRAY_SRV Texture1DArray;  
        D3D11_TEX2D_SRV      Texture2D;  
        D3D11_TEX2D_ARRAY_SRV Texture2DArray;  
        D3D11_TEX2DMS_SRV     Texture2DMS;  
        D3D11_TEX2DMS_ARRAY_SRV Texture2DMSArray;  
        D3D11_TEX3D_SRV      Texture3D;  
        D3D11_TEXCUBE_SRV     TextureCube;  
        D3D11_TEXCUBE_ARRAY_SRV TextureCubeArray;  
        D3D11_BUFFEREX_SRV    BufferEx;  
    };  
} D3D11_SHADER_RESOURCE_VIEW_DESC;
```


- CreateDDSTextureFromFile() takes care of all this!
 - Loads texture arrays from single DDS
 - Creates resource and resource views
- DirectXTex: texAssemble command line tool
 - Generates 2D Arrays / Cubemaps from multiple images
 - Outputs a single DDS
 - Add the texAssemble project to your solution...
 - ... and call it in your ResourceGenerator

- TexAssemble requires explicit specification of all images
 - We'll provide you a batch file to convert complete folders
 - See assignment
- TexAssemble does not provide MipMap-Generation and compression
 - Call texconv on the resulting DDS!

- Texture object declaration

```
Texture2DArray g_SpriteTex;
```

HLSL

- Binding shader resource variable (as usual)

```
m_pEffect->GetVariableByName("g_SpriteTex")->AsShaderResource()  
->SetResource(...);
```

C++

- Sampling the array

```
//position.xy: texture coordinates  
//position.z : array index  
float3 position;  
g_SpriteTex.Sample(samAnisotropic, position);
```

HLSL

- Array index position.z: rounded to next integer (array index)
- Use GetDimensions() to get the texture size in the shader

```
float3 dims; //(x-width, y-height, arraysize)  
g_SpriteTex.GetDimensions(dims.x, dims.y, dims.z); //dims is written here!
```

HLSL

- How to handle multiple sprite texture indices?
 - Only a single draw call for all particles (sorting!)
 - Use an array with a reasonable upper bound

```
Texture2DArray g_SpriteTex[5];
```

- Binding resources, two alternatives:

```
m_pEffect->GetVariableByName("g_SpriteTex")->AsShaderResource()  
->SetResourceArray(...);
```

```
m_pEffect->GetVariableByName("g_SpriteTex")->GetElement(index)  
->AsShaderResource()->SetResource (...);
```

- HLSL does not support texture fetches with dynamic indices
 - `switch()` with all possible indices does the job...
 - Won't win a price for the most beautiful code, though

- Same procedure as for your projectiles
- Add each explosion to a global list
 - Store texture index, position and current animation state `t`
 - Each frame, advance the animation state of all explosions
 - Use `fElapsedTime` and the total duration
 - Don't forget to delete explosions!
- On render, add all your explosions to the list of `SpriteVertex`
 - Sort and render along with your projectiles

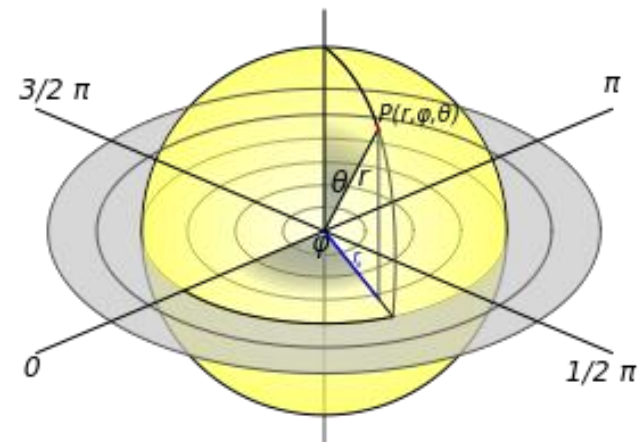
- Adding some fancyness: additional particles



- Create a number of „explosion particles“ at the center
 - Again, like your projectiles!
 - Update (speed, position, animation?, alpha?), render
- Storage
 - Either use a new global container for this
 - ... or unify all properties (explosion, particles, projectiles)
 - Using seperate containers might be better for experimenting though
- Initialization
 - Random direction
 - Random speed, size & position?

- Some hints for random directions
 - See the slides for Assignment 2 for random number generators and distributions!
 - Randomizing x, y, z in $[-1, 1]$ will result in a bad distribution
 - Alternative 1: **Rejection sampling**, discard if $\text{length}((x, y, z)) > 1$
 - Alternative 2: **Spherical coordinates**, two random angles θ in $[0^\circ; 180^\circ]$ and φ in $[0^\circ; 360^\circ]$:

$$\begin{pmatrix} x \\ y \\ z \end{pmatrix} = \begin{pmatrix} \sin(\theta) \cdot \cos(\varphi) \\ \sin(\theta) \cdot \sin(\varphi) \\ \cos(\theta) \end{pmatrix}$$



Questions?

(we're not done, don't run away yet!)

Maybe it's time to address
your particle system
performance.

At least a tiny bit.

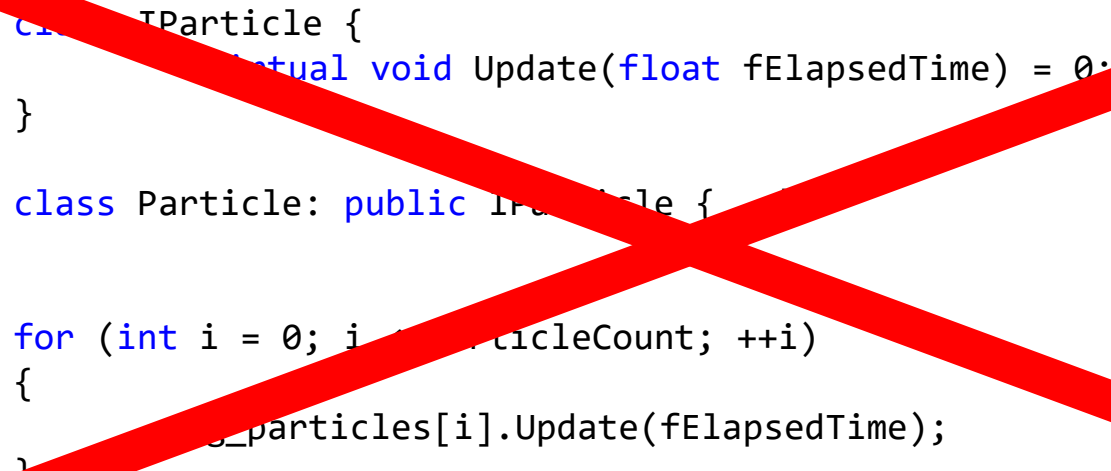
Study at home!



- Keep an eye on your function calls
 - Function calls aren't for free
 - Calling overridden functions is more expensive!
- Worst case scenario:

```
class IParticle {  
    virtual void Update(float fElapsedTime) = 0;  
}  
  
class Particle: public IParticle { ... }  
  
for (int i = 0; i < particleCount; ++i)  
{  
    g_particles[i].Update(fElapsedTime);  
}
```

- Keep an eye on your function calls
 - Function calls aren't for free
 - Calling overridden functions is more expensive!
 - Worst case scenario:



```
class IParticle {  
    virtual void Update(float fElapsedTime) = 0;  
}  
  
class Particle: public IParticle {  
  
    for (int i = 0; i < particleCount; ++i)  
    {  
        _particles[i].Update(fElapsedTime);  
    }  
}
```

- Keep an eye on your function calls
 - Instead:


```
struct Particle {  
    XMFLOAT3 m_position  
    // Speed, etc  
}  
  
for (int i = 0; i < particleCount; ++i)  
{  
    g_particles[i].m_position = ...;  
}
```

`std::vector` VS `std::list`


- Array of elements
- Consecutively stored in memory
- Single pointer to first element



- Each elements points to next & prev
- Can be anywhere in memory



- `std::list`
 - Removal at a random position is fast
 - Iteration over all elements is „okay-ish“
 - No random access – always iterate over all elements
 - Adding elements at random positions is okay, but memory allocations / deallocations can kill you



- `std::vector`
 - Iteration over all elements is fast
 - Random access is fast
 - Appending elements is fast - most of the time!
 - All elements are copied when no more space is available
 - Amount of reserved memory is doubled -> minimizes reallocation operations
 - Removal of the last element is fast
 - Removal of intermediate elements (`erase()`) is deadly...

- `std::vector`
 - Iteration over all elements is fast
 - Random access is fast
 - Appending elements is fast - most of the time!
 - All elements are copied when no more space is available
 - Amount of reserved memory is doubled -> minimizes reallocation operations
 - Removal of the last element is fast
 - Removal of intermediate elements (`erase()`) is deadly...



- Avoid removal of intermediate elements in `std::vector`
 - For „bulk removal“: instead of removing an element, „cut and paste“ the last one

```
int count = vec.size();
for (int i = 0; i < count; ++i) {

    bool remove = ...;
    if (remove) {
        --count;
        vec[i] = vec[count];
        --i;          // Check the same index again
    }

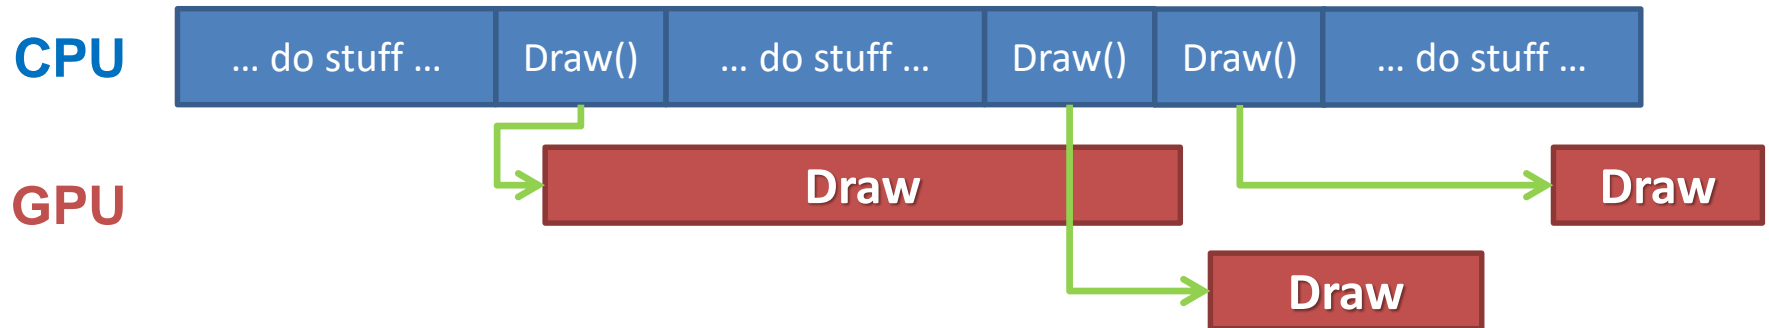
}


vec.resize(count);
```


- `std::map` and `std::unordered_map`
 - Use those if you need a „dictionary“ style container
 - `std::map`
 - (Binary) tree
 - Elements have an order
 - Access time: $O(\log n)$
 - `std::unordered_map`
 - Hash map
 - Elements have no order
 - Access time: „Usually“ constant, $O(1)$


- `std::vector` is the best choice most of the time
 - Keep in mind how the data is organized
 - Avoid `erase()`
 - Reserved memory isn't released if you resize a vector
 - So appending elements is basically free at some point
 - Only if your vector isn't a local variable!
- Always iterate „front to back“ over your data
 - Access the data in the order it's stored in memory
 - Keyword „caching“: The next values might already be stored in the fast CPU caches

- Keep in mind that data is processed in **parallel** and **asynchronously** on the GPU
 - Parallelism: enough data must be available to process
 - Asynchronism: synchronization must be minimized



- Keep the GPU busy!
 - Reduce the number of draw calls
 - In fact: Increase the data per draw call to avoid GPU idle time
 - Avoid synchronization points
 - GPU output is CPU input 
 - Could you use the data from the previous frame?
 - Could you draw something else during the idle-time?
 - Sidenote: Of course this cannot be completely avoided...



- Keep the GPU busy!
 - Reduce the number of draw calls
 - In fact: Increase the data per draw call to avoid GPU idle time
 - Avoid synchronization points
 - GPU output is CPU input 
 - Could you use the data from the previous frame?
 - Could you draw something else during the idle-time?
 - Sidenote: Of course this cannot be completely avoided...
 - Avoid overdraw
 - Many large, transparent particles are expensive

- Avoid unnecessary calculations in the shader:
Precalculate on the CPU
 - Combine transformation matrices

```
T3dVertexPSIn MeshVS(T3dVertexVSIn Input)
{
    T3dVertexPSIn Output = (T3dVertexPSIn) 0;
    Output.Pos = mul(float4(Input.Pos, 1.f), g_matObject);
    Output.Pos = mul(Output.Pos, g_matAnim);
    Output.Pos = mul(Output.Pos, g_matWorld);
    Output.Pos = mul(Output.Pos, g_matView);
    Output.Pos = mul(Output.Pos, g_matProjection);

    return Output;
}
```



```
T3dVertexPSIn MeshVS(T3dVertexVSIn Input)
{
    T3dVertexPSIn Output = (T3dVertexPSIn) 0;
    Output.Pos = mul(float4(Input.Pos, 1.f), g_WorldViewProjection);
    return Output;
}
```

- Precalculate per-frame constant factors in complex formulas

- Shaders are executed on SIMD Units
 - Single instruction, multiple data
 - Consequence: „Multiple data“ must be available to run the same instruction on...
- **Coherence** is important: Works best if all data is processed the same way
- Each group is as slow as its slowest member...
 - Avoid divergent branches
 - Avoid pixel shaders with highly varying local complexity

Always check your performance without a debugger!

- Build in Release configuration
- Start with **Ctrl + F5**
- `std` containers are quite slow in Debug build

Questions?